

Displaying Data from Multiple Tables

Objectives

- ▶ After completing this lesson, you should be able to do the following:
 - Write `SELECT` statements to access data from more than one table using equijoins and nonequijoins
 - Join a table to itself by using a self-join
 - View data that generally does not meet a join condition by using outer joins
 - Generate a Cartesian product of all rows from two or more tables

Tables

- ▶ **EMPLOYEES** (`EMPLOYEE_ID` , `FIRST_NAME` , `LAST_NAME` , `EMAIL` , `PHONE_NUMBER` , `HIRE_DATE` , `JOB_ID` , `SALARY` , `COMMISSION_PCT` , `MANAGER_ID` , `DEPARTMENT_ID`)
- ▶ **DEPARTMENTS** (`DEPARTMENT_ID` , `DEPARTMENT_NAME` , `MANAGER_ID`,`LOCATION_ID`)
- ▶ **LOCATIONS** (`LOCATION_ID`,`STREET_ADDRESS`,`POSTAL_CODE`,`CITY`,`STATE_PROVINCE`,`COUNTRY_ID`)

- ▶ A *join* is used to view information from multiple tables.
- ▶ You can *join* tables together to view information from more than one table.

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID
1	200 Watson	20
2	201 Hartstein	20
3	202 Ky	20
4	203 Higgins	110
...
10	174 Deel	80
19	176 Deier	80
20	177 Deier	80

DEPARTMENT_ID	DEPARTMENT_NAME	LOCATION_ID
1	Administration	1700
2	Marketing	1800
3	Shipping	1500
4	IT	1400
5	Sales	2500
6	Executive	1700
7	Accounting	1700
8	Contracting	1700

EMPLOYEE_ID	DEPARTMENT_ID	DEPARTMENT_NAME
1	100	90 Executive
2	100	90 Executive
...
17	200	20 Marketing
18	205	110 Accounting
19	206	110 Accounting

Types of Joins

- ▶ Joins that are compliant with the SQL:1999 standard include the following:
 - Cross joins
 - Natural joins
 - USING clause
 - Full (or two-sided) outer joins
 - Arbitrary join conditions for outer joins

Joining Tables Using SQL:1999 Syntax

- ▶ Use a join to query data from more than one table:

```
SELECT table1.column, table2.column
FROM table1
[NATURAL JOIN table2] |
[JOIN table2 USING (column_name)] |
[JOIN table2
ON (table1.column_name = table2.column_name)] |
[LEFT|RIGHT|FULL OUTER JOIN table2
ON (table1.column_name = table2.column_name)] |
[CROSS JOIN table2];
```

`table1.column` denotes the table and column from which data is retrieved

NATURAL JOIN joins two tables based on the same column name

JOIN table USING column_name performs an equijoin based on the column name

JOIN table ON table1.column_name performs an equijoin based on the condition in the **ON** clause, `= table2.column_name`

LEFT/RIGHT/FULL OUTER is used to perform outer joins

CROSS JOIN returns a Cartesian product from the two tables

Creating Natural Joins

- ▶ You can join tables automatically based on columns in the two tables that have matching data types and names. You do this by using the keywords **NATURAL JOIN**.
 - It selects rows from the two tables that have equal values in all matched columns.
 - If the columns having the same names have different data types, an error is returned.

Retrieving Records with Natural Joins

```
SELECT department_id, department_name,
       location_id, city
FROM   departments
NATURAL JOIN locations; J1
```

DEPARTMENT_ID	DEPARTMENT_NAME	LOCATION_ID	CITY
1	60 IT	1400	Southlake
2	50 Shipping	1500	South San Francisco
3	10 Administration	1700	Seattle
4	90 Executive	1700	Seattle
5	110 Accounting	1700	Seattle
6	190 Contracting	1700	Seattle
7	20 Marketing	1800	Toronto
8	80 Sales	2500	Oxford

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
30	Purchasing	114	1700
40	Human Resources	203	2400
50	Shipping	121	1500
60	IT	103	1400
70	Public Relations	204	2700
80	Sales	145	2500
90	Executive	100	1700
100	Finance	105	1700

LOCATION_ID	STREET_ADDRESS	POSTAL_CODE	CITY	STATE_PROVINCE	COUNTRY_ID
1000	1201 Via Cola di Rie	00089	Rome	-	IT
1100	60001 Calle della Vite	10034	Venice	-	IT
1200	2011 Shimbashi	10000	Tokyo	Tokyo Prefecture	JP
1300	9405 Kamoyu-cho	40021	Wakuhime	-	JP
1400	2014 Lakeshore Blvd	20102	Southlake	-	CA
1500	2011 Interiors Blvd	90236	South San Francisco	California	US
1600	2007 Zappala St	03080	South Brunswick	New Jersey	US
1700	2004 Charney Rd	98109	Seattle	Washington	US
1800	147 Spadina Ave	M5V 1L7	Toronto	Ontario	CA
1900	6092 Bonwood St	V5H 4T2	Whitehorse	Yukon	CA

Natural by using a WHERE clause

- ▶ limits the rows of output to those with a department ID equal to 20 or 50:

```
SELECT department_id, department_name,
       location_id, city
FROM   departments
NATURAL JOIN locations
WHERE  department_id IN (20, 50); J2
```

DEPARTMENT_ID	DEPARTMENT_NAME	LOCATION_ID	CITY
20	Marketing	1800	Toronto
50	Shipping	1500	South San Francisco

2 rows returned in 0.03 seconds [CSV Export](#)

Creating Joins with the USING Clause

- ▶ Natural joins use all columns with matching names and data types to join the tables.
- ▶ If several columns have the same names but the data types do not match, natural join can be applied by using the **USING** clause to specify the columns that should be used for an equijoin.
- ▶ Use the **USING** clause to match only one column when more than one column matches.
- ▶ Do not use a table name or alias in the referenced columns.
- ▶ The **NATURAL JOIN** and **USING** clauses are mutually exclusive.

```
SELECT l.city, d.department_name
FROM   locations l JOIN departments d USING (location_id)
WHERE  location_id = 1400; J3
```

- ▶ The following statement is invalid because the **LOCATION_ID** is qualified in the **WHERE** clause:

```
SELECT l.city, d.department_name
FROM   locations l JOIN departments d USING (location_id)
WHERE  d.location_id = 1400; J4
```

- ▶ **Note:** The same restriction also applies to **NATURAL** joins. Therefore, columns that have the same name in both tables must be used without any qualifiers.

Joining Column Names

EMPLOYEES		DEPARTMENTS	
EMPLOYEE_ID	DEPARTMENT_ID	DEPARTMENT_ID	DEPARTMENT_NAME
1	200	10	Administration
2	201	20	Marketing
3	202	20	Marketing
4	205	110	Shipping
5	206	110	Shipping
6	100	90	Shipping
7	101	90	Shipping
8	102	90	Shipping
9	103	60	IT
10	104	60	IT
11	107	60	IT
12	124	90	Sales
13	143	90	Sales
14	142	90	Sales
15	143	90	Sales
16	144	90	Sales
17	149	90	Sales
18	174	90	Sales
19	176	90	Sales
20	178	90	Sales

Foreign key Primary key

- ▶ The relationship between the EMPLOYEES and DEPARTMENTS tables is an *equijoin*;
- ▶ **Note:** Equijoins are also called *simple joins* or *inner joins*.

Retrieving Records with the USING Clause

- ▶ Ex. joins the DEPARTMENT_ID column in the EMPLOYEES and DEPARTMENTS tables, and thus shows the location where an employee works.

```
SELECT employees.employee_id, employees.last_name,
       departments.location_id, department_id
FROM   employees JOIN departments
       USING (department_id); J5
```

EMPLOYEE_ID	LAST_NAME	LOCATION_ID	DEPARTMENT_ID
1	200 Whalen	1700	10
2	201 Hartstein	1800	20
3	202 Fay	1800	20
4	205 Higgins	1700	110
5	206 Gietz	1700	110
6	100 King	1700	90
7	101 Kochhar	1700	90
8	102 De Haan	1700	90
9	103 Hunold	1400	60
10	104 Ernst	1400	60

Qualifying Ambiguous Column Names

- ▶ You need to qualify the names of the columns with the table name to avoid ambiguity.
- ▶ Without the table prefixes, the DEPARTMENT_ID column in the SELECT list could be from either the DEPARTMENTS table or the EMPLOYEES table.
- ▶ It is necessary to add the table prefix to execute your query

```
SELECT employees.employee_id, employees.last_name,
       departments.department_id, departments.location_id
FROM   employees JOIN departments
ON     employees.department_id = departments.department_id; J6
```

- ▶ If there are no common column names between the two tables, there is no need to qualify the columns. However, using the table prefix improves performance, because you tell the Oracle server exactly where to find the columns.

- ▶ **Note:** When joining with the USING clause, you cannot qualify a column that is used in the USING clause itself. Furthermore, if that column is used anywhere in the SQL statement, you cannot alias it.

Using Table Aliases

- ▶ Qualifying column names with table names can be very time consuming, particularly if table names are lengthy.
- ▶ You can use *table aliases* instead of table names.
- ▶ Just as a column alias gives a column another name, a table alias gives a table another name.
- ▶ How table aliases are identified in the FROM clause
- ▶ The table name is specified in full, followed by a space and then the table alias.

```
SELECT e.employee_id, e.last_name,
       d.location_id, department_id
FROM   employees e JOIN departments d
       USING (department_id); J7
```

Guidelines

- Table aliases can be up to 30 characters in length, but shorter aliases are better than longer ones.
- If a table alias is used for a particular table name in the FROM clause, then that table alias must be substituted for the table name throughout the SELECT statement.
- Table aliases should be meaningful.
- The table alias is valid for only the current SELECT statement.

Creating Joins with the ON Clause

- ▶ The join condition for the natural join is basically an equijoin of all columns with the same name.
- ▶ Use the ON clause to specify arbitrary conditions or specify columns to join.
- ▶ The join condition is separated from other search conditions.
- ▶ The ON clause makes code easy to understand.

Retrieving Records with the ON Clause

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
FROM   employees e JOIN departments d
ON     (e.department_id = d.department_id); J8
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID	DEPARTMENT_ID_1	LOCATION_ID
1	200 Whalen	10	10	1700
2	201 Hartstein	20	20	1800
3	202 Fay	20	20	1800
4	205 Higgins	110	110	1700
5	206 Gietz	110	110	1700
6	100 King	90	90	1700
7	101 Kochhar	90	90	1700
8	102 De Haan	90	90	1700
9	103 Hunold	60	60	1400
10	104 Ernst	60	60	1400

- ▶ Note: You can also use the ON clause to join columns that have different names.

Self-Joins Using the ON Clause

- ▶ Sometimes you need to join a table to itself.
- Ex. To find the name of each employee's manager

EMPLOYEES (WORKER)				EMPLOYEES (MANAGER)			
EMPLOYEE_ID	LAST_NAME	MANAGER_ID		EMPLOYEE_ID	LAST_NAME		
1	100 King	(null)		1	100 King		
2	101 Kochhar	100		2	101 Kochhar		
3	102 De Haan	100		3	102 De Haan		
4	103 Hunold	102		4	103 Hunold		
5	104 Ernst	103		5	104 Ernst		
6	107 Lorentz	103		6	107 Lorentz		
7	124 Mourgos	100		7	124 Mourgos		
8	141 Rajis	124		8	141 Rajis		
9	142 Davies	124		9	142 Davies		
10	143 Matos	124		10	143 Matos		

MANAGER_ID in the WORKER table is equal to
EMPLOYEE_ID in the MANAGER table.

Self-Joins Using the ON Clause

- ▶ The ON clause can also be used to join columns that have different names, within the same table or in a different table.

```
SELECT e.last_name emp, m.last_name mgr
FROM   employees e JOIN employees m
ON     (e.manager_id = m.employee_id); J9
```

EMP	MGR
1 Abel	Zlotkey
2 Davies	Mourgos
3 De Haan	King
4 Ernst	Hunold
5 Fay	Hartstein
6 Gietz	Higgins
7 Grant	Zlotkey
8 Hartstein	King

Applying Additional Conditions to a Join

- ▶ You can apply additional conditions to the join.
- ▶ Ex. Performs a join on the EMPLOYEES and DEPARTMENTS tables and, in addition, displays only employees who have a manager ID of 149.

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
FROM   employees e JOIN departments d
ON     (e.department_id = d.department_id)
AND    e.manager_id = 149; J10
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID	DEPARTMENT_ID_1	LOCATION_ID
1	174 Abel	80	80	2500
2	176 Taylor	80	80	2500

- ▶ To add additional conditions to the ON clause, you can add AND clauses.

- ▶ Alternatively, you can use a WHERE clause to apply additional conditions:

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
FROM   employees e JOIN departments d
ON     (e.department_id = d.department_id)
WHERE  e.manager_id = 149; J11
```

Creating Three-Way Joins with the ON Clause

- ▶ A three-way join is a join of three tables.
- ▶ In SQL:1999-compliant syntax, joins are performed from left to right. So, the first join to be performed is `EMPLOYEES JOIN DEPARTMENTS`.
- ▶ The first join condition can reference columns in `EMPLOYEES` and `DEPARTMENTS` but cannot reference columns in `LOCATIONS`.
- ▶ The second join condition can reference columns from all three tables.

```
SELECT employee_id, city, department_name
FROM employees e
JOIN departments d
ON d.department_id = e.department_id
JOIN locations l
ON d.location_id = l.location_id; J12
```

EMPLOYEE_ID	CITY	DEPARTMENT_NAME
1	100 Seattle	Executive
2	101 Seattle	Executive
3	102 Seattle	Executive
4	103 Southlake	IT
5	104 Southlake	IT
6	107 Southlake	IT
7	124 South San Francisco	Shipping
8	141 South San Francisco	Shipping

Nonequijoins

- ▶ A nonequijoin is a join condition containing something other than an equality operator.
- ▶ The relationship between the `EMPLOYEES` table and the `JOB_GRADES` table is an example of a nonequijoin
- ▶ The relationship is obtained using an operator other than equality (=).

EMPLOYEES

LAST_NAME	SALARY
1 Whalen	4400
2 Hartstein	13000
3 Fay	6000
4 Higgins	12000
5 Gietz	8300
6 King	24000
7 Kochhar	17000
8 De Haan	17000
9 Hunold	9000
10 Ernst	6000

JOB_GRADES

GRADE_LEVEL	LOWEST_SAL	HIGHEST_SAL
1 A	1000	2999
2 B	3000	5999
3 C	6000	9999
4 D	10000	14999
5 E	15000	24999
6 F	25000	40000

Salary in the `EMPLOYEES` table must be between lowest salary and highest salary in the `JOB_GRADES` table.

Retrieving Records with Nonequijoins

Ex. creates a nonequijoin to evaluate an employee's job title. The salary must be *between* any pair of the low and high salary ranges.

```
SELECT e.last_name, e.salary, j.job_title
FROM employees e JOIN jobs j
ON e.salary
BETWEEN j.min_salary AND j.max_salary; J13
```

LAST_NAME	SALARY	GRADE_LEVEL
1 Vargas	2500	A
2 Matos	2600	A
3 Davies	3100	B
4 Rajs	3500	B
5 Lorentz	4200	B
6 Whalen	4400	B
7 Mourgos	5800	B
8 Ernst	6000	C

Outer Joins

- ▶ If a row does not satisfy a join condition, the row does not appear in the query result.

DEPARTMENTS

DEPARTMENT_NAME	DEPARTMENT_ID
1 Administration	10
2 Marketing	20
3 Shipping	50
4 IT	60
5 Sales	80
6 Executive	90
7 Accounting	110
8 Contracting	190

EMPLOYEES

DEPARTMENT_ID	LAST_NAME
1	10 Whalen
2	20 Hartstein
3	20 Fay
4	110 Higgins
5	110 Gietz
6	90 King
7	90 Kochhar
8	90 De Haan
9	60 Hunold
10	60 Ernst

There are no employees in department 190.

INNER Versus OUTER Joins

- ▶ In SQL:1999, the join of two tables returning only matched rows is called an inner join.
- ▶ A join between two tables that returns the results of the inner join as well as the unmatched rows from the left (or right) tables is called a left (or right) outer join.
- ▶ A join between two tables that returns the results of an inner join as well as the results of a left and right join is a full outer join.
 - There are three types of outer joins:
 - LEFT OUTER
 - RIGHT OUTER
 - FULL OUTER

LEFT OUTER JOIN

- Query retrieves all rows in the `EMPLOYEES` table, which is the table on the left even if there is no match in the `DEPARTMENTS` table.

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e LEFT OUTER JOIN departments d
ON (e.department_id = d.department_id) ; J14
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID	DEPARTMENT_NAME
1	Whalen	10	Administration
2	Hartstein	20	Marketing
3	Fay	20	Marketing
4	Higgins	110	Accounting
...			
18	Abel	80	Sales
19	Taylor	80	Sales
20	Grant	(null)	(null)

RIGHT OUTER JOIN

- Query retrieves all rows in the `DEPARTMENTS` table, which is the table on the right even if there is no match in the `EMPLOYEES` table.

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e RIGHT OUTER JOIN departments d
ON (e.department_id = d.department_id) ; J15
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID	DEPARTMENT_NAME
1	Whalen	10	Administration
2	Hartstein	20	Marketing
3	Fay	20	Marketing
...			
18	Higgins	110	Accounting
19	Gietz	110	Accounting
20	(null)	(null)	Contracting

FULL OUTER JOIN

- Query retrieves all rows in the `EMPLOYEES` table, even if there is no match in the `DEPARTMENTS` table. It also retrieves all rows in the `DEPARTMENTS` table, even if there is no match in the `EMPLOYEES` table.

```
SELECT e.last_name, d.department_id, d.department_name
FROM employees e FULL OUTER JOIN departments d
ON (e.department_id = d.department_id) ; J16
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID	DEPARTMENT_NAME
1	Whalen	10	Administration
2	Hartstein	20	Marketing
3	Fay	20	Marketing
...			
18	Abel	80	Sales
19	Taylor	80	Sales
20	Grant	(null)	(null)
21	(null)	190	Contracting

Cartesian Products

- When a join condition is invalid or omitted completely, the result is a *Cartesian product*, in which all combinations of rows are displayed.
- All rows in the first table are joined to all rows in the second table.
- A Cartesian product tends to generate a large number of rows, and the result is rarely useful.
- You should always include a valid join condition unless you have a specific need to combine all rows from all tables.

Generating a Cartesian Product

- shows employee last name and department name from the `EMPLOYEES` and `DEPARTMENTS` tables.
- Because no join condition has been specified

EMPLOYEES (20 rows)

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID
1	200 Whalen	10
2	201 Hartstein	20
...		
19	176 Taylor	80
20	178 Grant	(null)

DEPARTMENTS (8 rows)

DEPARTMENT_ID	DEPARTMENT_NAME	LOCATION_ID
1	10 Administration	1700
2	20 Marketing	1800
3	50 Shipping	1500
4	60 IT	1400
5	80 Sales	2500
6	90 Executive	1700
7	110 Accounting	1700
8	190 Contracting	1700

Cartesian
product:
20 x 8 = 160
rows

EMPLOYEE_ID	DEPARTMENT_ID	LOCATION_ID
1	100	10
2	101	10
...		
158	200	190
157	201	190
158	202	190
159	205	190
160	206	190

```
SELECT last_name, department_name
FROM employees
CROSS JOIN departments ; J17
```

EMPLOYEE_ID	LAST_NAME	DEPARTMENT_NAME
1	Abel	Administration
2	Davies	Administration
3	De Haan	Administration
4	Ernst	Administration
...		
159	Whalen	Contracting
160	Zlotkey	Contracting